



Embedded Systems
Research Group

Computer Architecture

Caches

Filipe Salgado, PhD Student @ ESRG

R&D Centre ALGORITMI
Department of Industrial Electronics
School of Engineering,
University of Minho, Guimarães - PORTUGAL





Cache

- Overview
- Introduction
- What to cache?
- Logical Organization
- Cachegen
- Cache implementation "from scratch"





Overview

- Introduction
 - Why to Cache?
 - Memory Hierarchy
- What to cache?
 - Locality of Reference
- Logical Organization
 - Concepts, words, lines, sets, tags
 - Types of Caches (organizational models)





Overview

- Cachegen
 - Caches generation
 - Generate RTL in Quartus II
- Cache implementation "from scratch"
 - Direct mapped cache
 - Metrics and projection
 - Storage structures
 - test-bench





Cache

- Overview
- **Introduction**
- What to cache?
- Logical Organization
- Cachegen
- Cache implementation "from scratch"





Why to cache?

- Why to cache?
 - Memory technology evolves differently from silicon technology
 - Different speeds:
 - DDR SDRAM: **100MHz-266MHz**
 - Pentium: **1.3GHz – 3.8GHz**
 - SDRAM requires **multiple cycles** to access content
 - Memory access is a performance bottleneck in typical computer architectures
 - Cache attenuates that performance difference
 - Storing memory content in a fast access (but area expensive) memory, closer to the processor core.





Memory Hierarchy

- Memory Hierarchy
 - Fast caches are very expensive!
 - In Power and area!!!
 - About 20% to 30% of chip Area!
 - So must have reduced storage capacity
 - Solution is to create slower caches with increased storage capacity!
 - Compose a memory hierarchy
 - Cache L1 – fastest, most expensive
 - Cache L2 – fast, less expensive, more storage capacity
 - ?Cache L3?
 - SDRam
 - Mass storage memory






Memory Hierarchy

- Memory Hierarchy

- Cache L1 – fastest, expensivest
- Cache L2 – fast, less expensive, more storage capacity
- Cache L3 - ...
- SDRam
- Mass storage memory

- Convenience Store
Example 😊





Cache



Embedded Systems
Research Group

- Overview
- Introduction
- **What to cache?**
- Logical Organization
- Caching
- Cache implementation "from scratch"





What to cache?

- Data caching is done according to straight information management policies
 - How to decide what should be cached?
- Locality of Reference:
 - Principle of Temporal Locality
 - Principle of Spatial Locality





Locality of Reference

- Principle of Temporal Locality
 - “If a particular memory location at one point is required, it is very likely that it will be necessary again in a near future.”
- Convenience store example ;)
 - If you go there to buy eggs and they don't have, when they go to the warehouse, they'll get more than one box.
- Usage ex.: Data memory (variables in loop cycles)





Locality of Reference

- Principle of Spatial Locality
 - “If a particular memory location at one point is required, it is very likely that the nearby locations will be referenced too in a near future.”
 - Convenience store example ;)
 - If you go there to buy potatoes and they are out of stock, when they go to the market, they’ll take advantage of the trip and also buy carrots, tomatoes, etc...
 - Usage ex.: Instructions stored in Code Memory





Cache



Embedded Systems
Research Group

- Overview
- Introduction
- What to cache?
- **Logical Organization**
- Cachegen
- Cache implementation "from scratch"

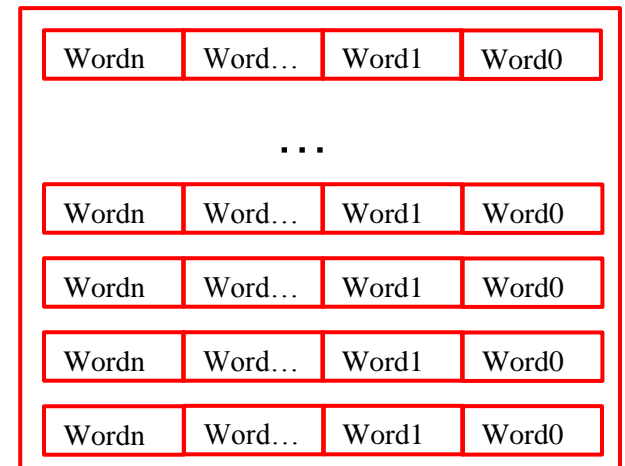
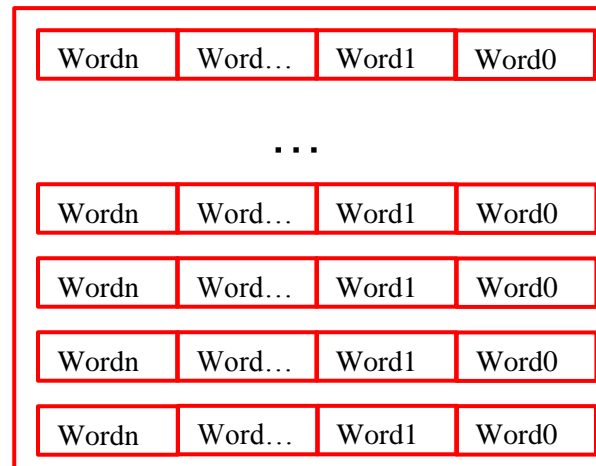




Logical Organization

- How data is organized in cache?
 - Words
 - Lines
 - Sets

...





Logical Organization

- How to know **where** is **what**?

(considering direct mapped and set-associative)

- Addresses are stored in a dedicated memory - **tag memory**
- Some part of the address is used to address the Tag memory – **line address**
- Another part of the address is used to address the word – **word address**

(further explanations ahead)





Types of cache



Embedded Systems
Research Group


- Some organizational models are used:
 - **Fully associative**
 - **Direct mapped**
 - **Set associative**
 - **Hyper associative**





Types of cache

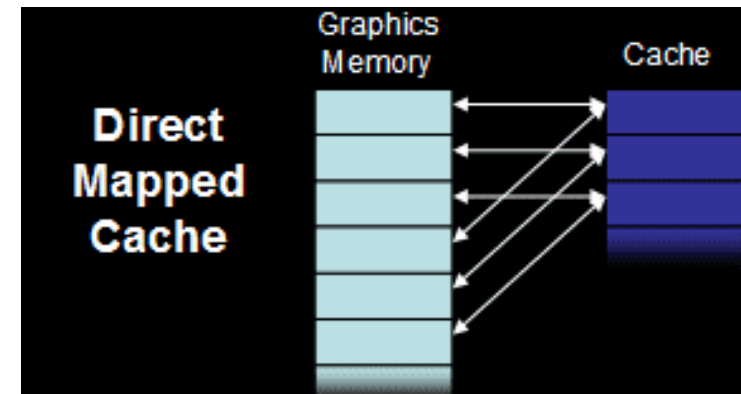
- Fully associative:
 - Data is stored in the cache sets according with the access order.
 - Any memory content can be anywhere in the set
 - Required a bigger tag memory
 - Uses a CAM (content addressable memory)
 - Reverse memory: user provides the content, it gives the address where it is stored
 - Address is stored in the CAM
 - After read the cam user knows if requested location is available or not.





Types of cache

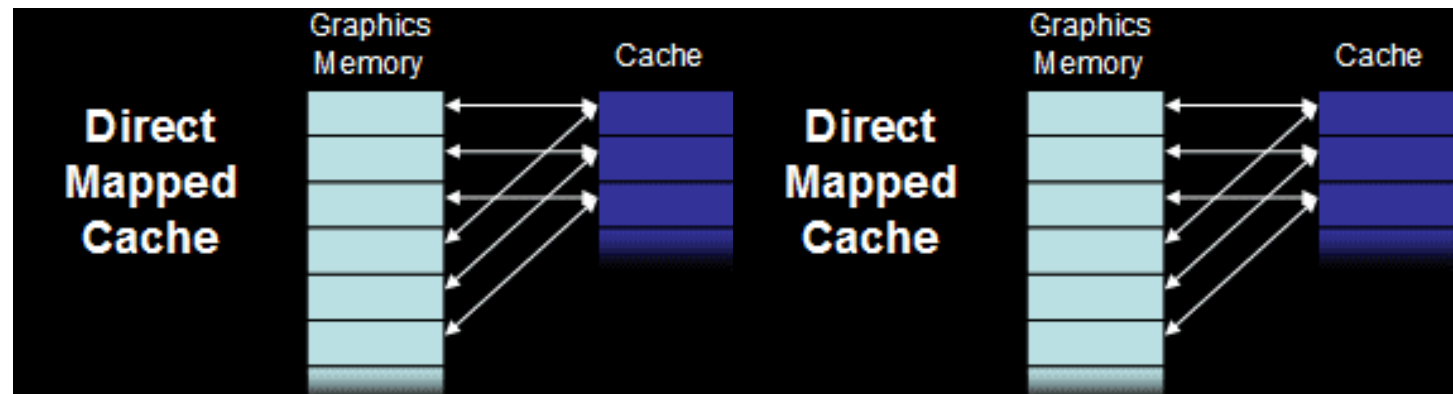
- Direct Mapped:
 - Easier to implement
 - Each cache line corresponds to a subset of memory positions
 - Smaller tag memory





Types of cache

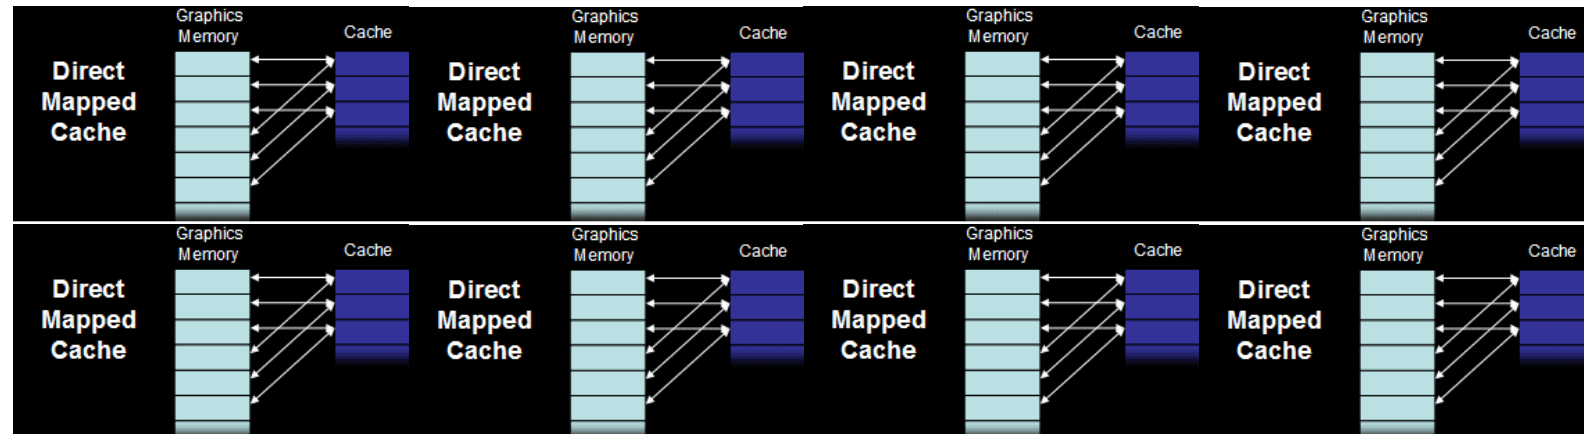
- N-way set associative (typically N is 2,4,8 sets):
 - Same principle of direct mapped cache
 - Associativity in powers of 2!
 - Implemented with the replications of groups of cache set+tag memory





Types of cache

- Hyper-associative (typically 16,32,64,128,... sets):
 - Its an implementation of a N-way set associative with a huge number of small sets
 - Common in super scalar implementations
 - Increased number of sets to increase hit rate in applications with several small data work sets





Cache



Embedded Systems
Research Group

- Overview
- Introduction
- What to cache?
- Logical Organization
- **Cachegen**
- Cache implementation "from scratch"





Cachegen

- It's a software developed by
 - Peter Yiannacouras for partial fulfillment of **BASc** at the University of Toronto
 - For Linux
- Simple cache generator, able to generate synthesizable **Verilog** code for **associative**, **direct mapped** and **2way set associative** caches.
- Generates only "one word per line" cache sets
- Useful for educational purposes!
 - Lets try it ;)





Cachegen

- Go to Linux
- Extract the tarball

Run the following commands:

```
gattuso@ubuntu:~$ cd ~/Desktop/cachegeen/
gattuso@ubuntu:~/Desktop/cachegeen$ ls
design  genall.sh  main.c  Makefile  README.txt
gattuso@ubuntu:~/Desktop/cachegeen$ make
gcc -g -o cachegeen main.c -lm
gattuso@ubuntu:~/Desktop/cachegeen$ ls
cachegeen  design  genall.sh  main.c  Makefile  README.txt
gattuso@ubuntu:~/Desktop/cachegeen$ ./cachegeen

CAM GENERATION
Usage: cachegeen -c <depth> <width> <output location>
ASSOCIATIVE CACHE GENERATION
Usage: cachegeen -a [-hei] <cache depth> <address width> <data width> <output location>
DIRECT MAPPED CACHE GENERATION
Usage: cachegeen -d [-hei] <cache depth> <address width> <data width> <output location>
TWO WAY SET ASSOCIATIVE CACHE GENERATION
Usage: cachegeen -t [-heiw] <cache depth> <address width> <data width> <output location>

OPTIONS:
-h Cache Hit - outputs a signal that indicates when a cache hit occurs
-e Expand Read - Adds an extra cycle to the read operation improving fmax
-i Intercept - passes read write signals to slave only when necessary
-w Expand Write - Adds an extra cycle to the write operation improving fmax
gattuso@ubuntu:~/Desktop/cachegeen$
```




Cachegen

- Cache depth?
- Address width?
- Data width?

```
gattuso@ubuntu:~$ cd ~/Desktop/cache/
gattuso@ubuntu:~/Desktop/cache$ ls
design  genall.sh  main.c  Makefile  README.txt
gattuso@ubuntu:~/Desktop/cache$ make
gcc -g -o cache main.c -lm
gattuso@ubuntu:~/Desktop/cache$ ls
cache  design  genall.sh  main.c  Makefile  README.txt
gattuso@ubuntu:~/Desktop/cache$ ./cache

CAM GENERATION
Usage: cache -c <depth> <width> <output location>
ASSOCIATIVE CACHE GENERATION
Usage: cache -a [-hei] <cache depth> <address width> <data width> <output location>
DIRECT MAPPED CACHE GENERATION
Usage: cache -d [-hei] <cache depth> <address width> <data width> <output location>
TWO WAY SET ASSOCIATIVE CACHE GENERATION
Usage: cache -t [-heiw] <cache depth> <address width> <data width> <output location>

OPTIONS:
-h Cache Hit - outputs a signal that indicates when a cache hit occurs
-e Expand Read - Adds an extra cycle to the read operation improving fmax
-i Intercept - passes read write signals to slave only when necessary
-w Expand Write - Adds an extra cycle to the write operation improving fmax
gattuso@ubuntu:~/Desktop/cache$
```




Cachegen



Embedded Systems
Research Group

- **Cache depth** - is the number of cache lines present in the cache set and tag memory. In other words is the number of **cache lines**
- **Address width** – is the number of bits available to address memory. It's a consequence of the architecture and relates with the addressable space.
- **Data width** – is the number of bits per cache line. In other words is the **line size**.





Cachegen

- Metrics examples:
 - cache dept 32
 - addr width (address space) 8
 - data width (line size) 16
- **What's the storage capacity of this cache?**
- **What's the addressable space?**

```
gattuso@ubuntu:~/Desktop/cache$ ./cache -a -h 32 8 16 ./associative
Creating directory ./associative
Generating 32x16 associative cache ... done
Generating 32:5 encoder ... done
Generating 32x8 cam ... done
gattuso@ubuntu:~/Desktop/cache$ ./cache -d -h 32 8 16 ./direct
Creating directory ./direct
Generating 32x16 direct-mapped cache ... done
gattuso@ubuntu:~/Desktop/cache$ ./cache -t -h 32 8 16 ./setassociative
Creating directory ./setassociative
Generating 32x16 two way set associative cache ... done
gattuso@ubuntu:~/Desktop/cache$ ls
associative  design  genall.sh  Makefile  setassociative
cache       direct  main.c    README.txt
gattuso@ubuntu:~/Desktop/cache$
```





- Back to windows
 - Copy the Direct mapped Design
 - Create a Vivado Project for Zybo board
 - Include the Verilog file
 - **Synthesise!**
- **Generate RTL View**

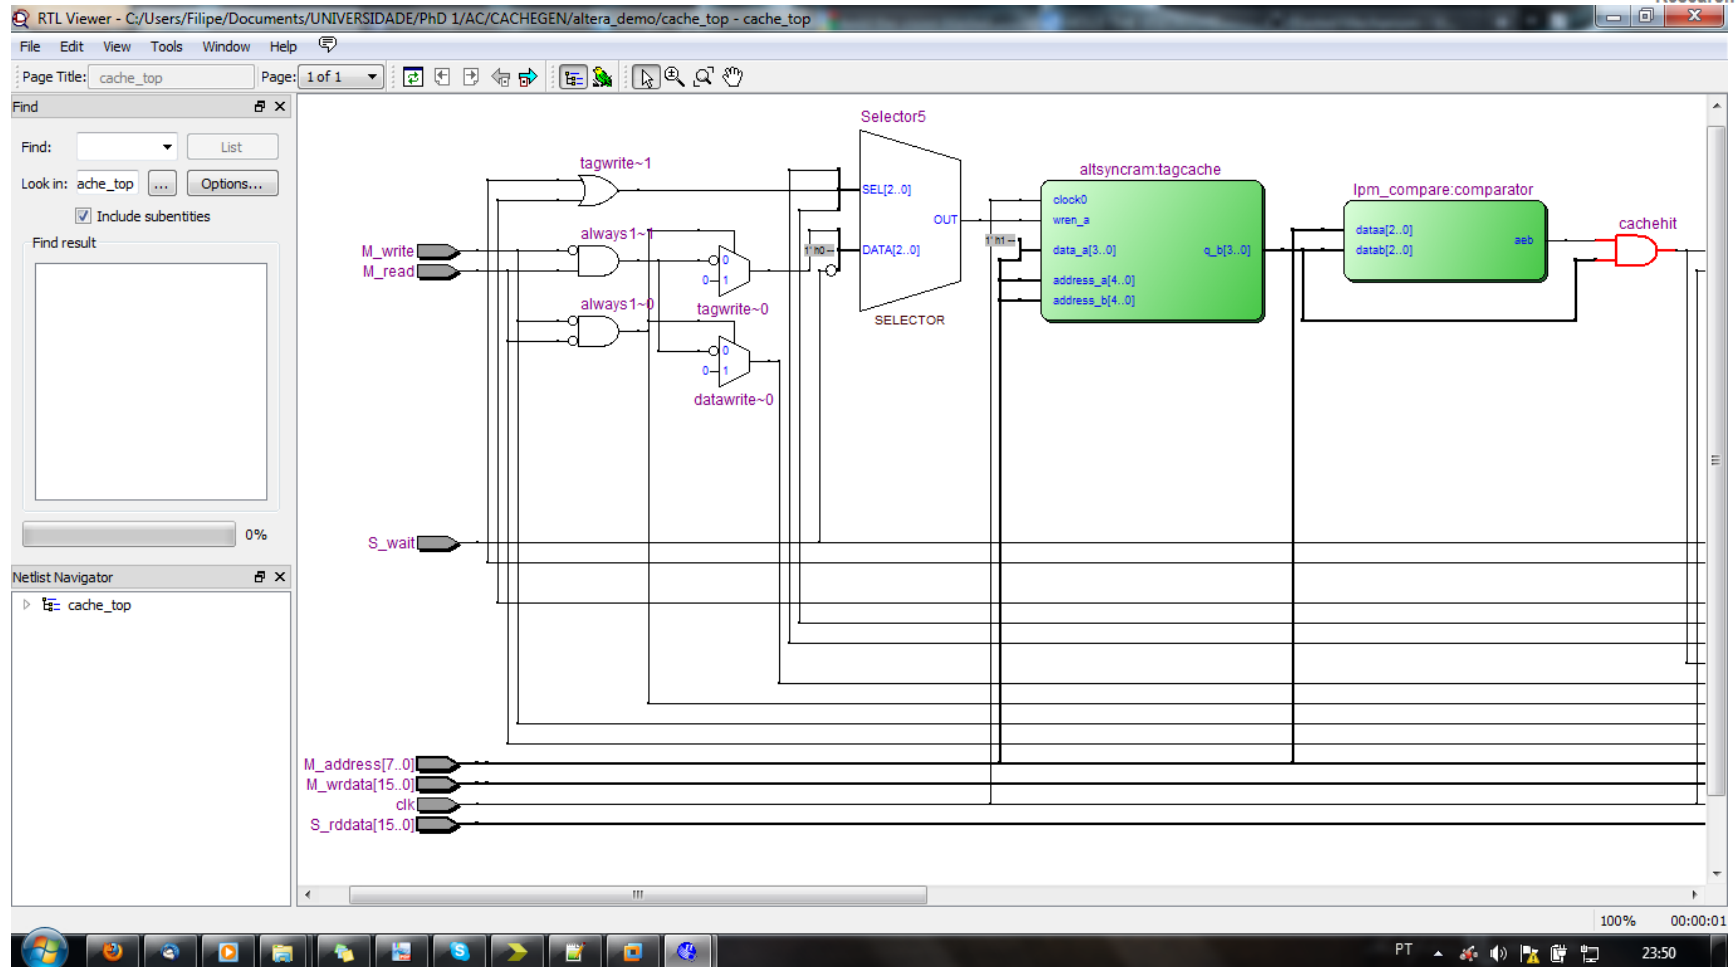




Cachegen



Embedded Systems
Research Group

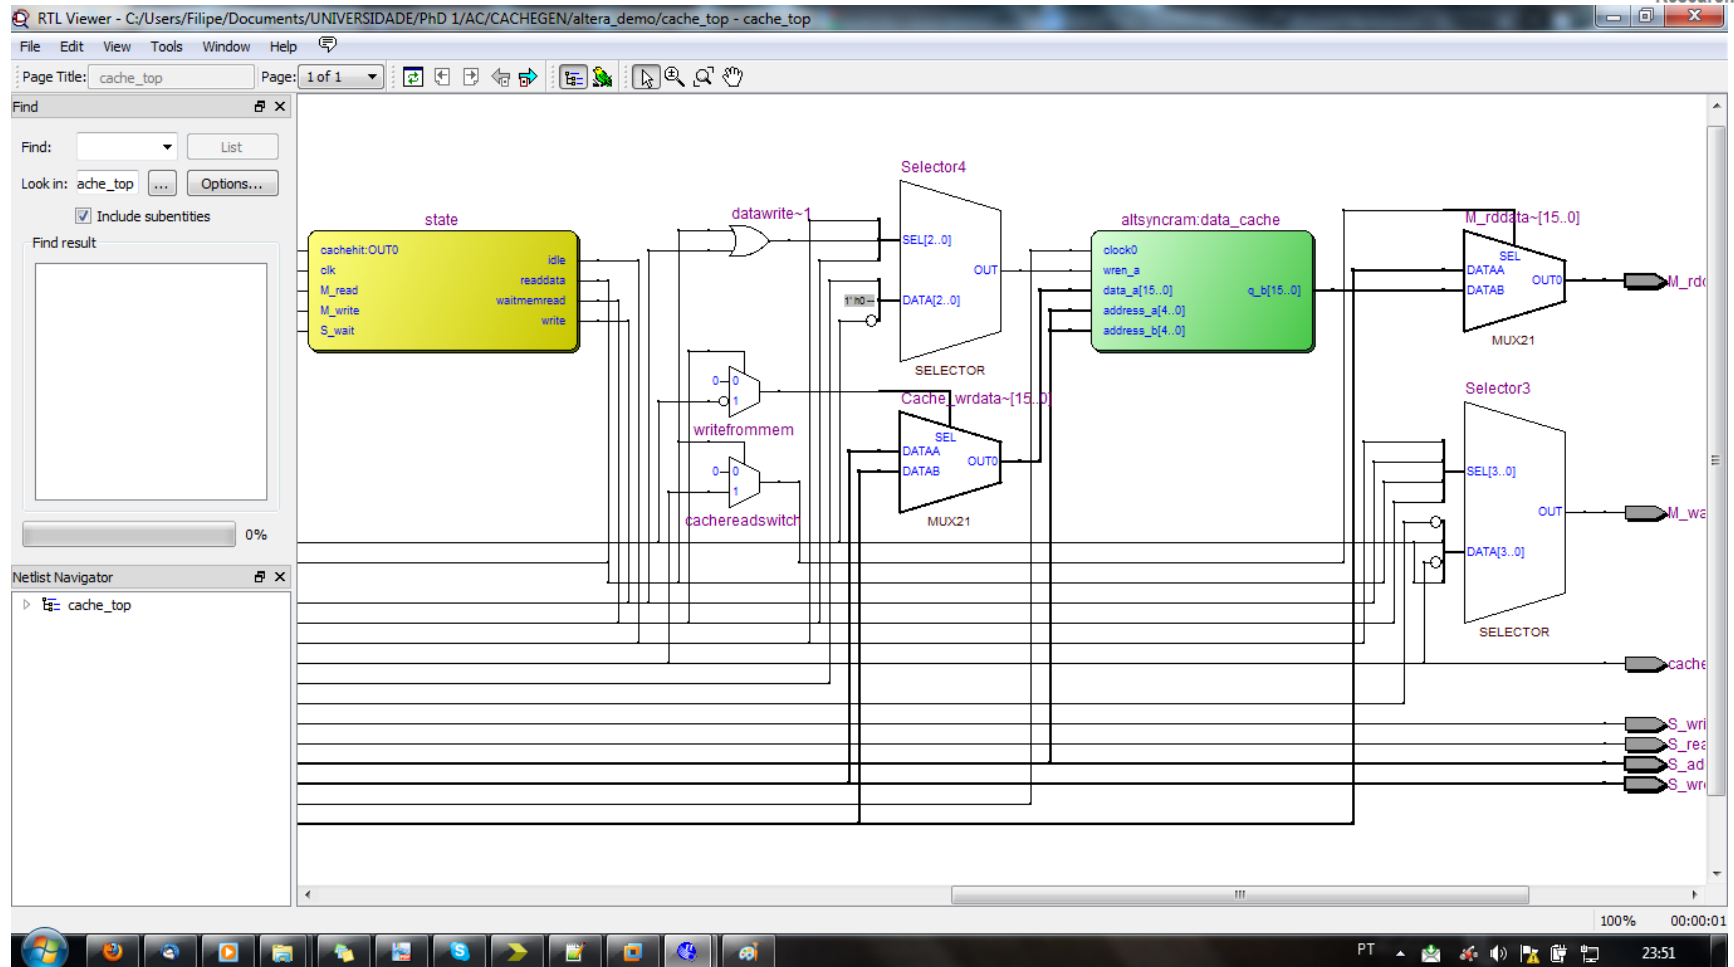




Cachegen



Embedded Systems
Research Group





Cache



Embedded Systems
Research Group

- Overview
- Introduction
- What to cache?
- Logical Organization
- Cachegen
- Cache implementation "from scratch"





Working principle basis

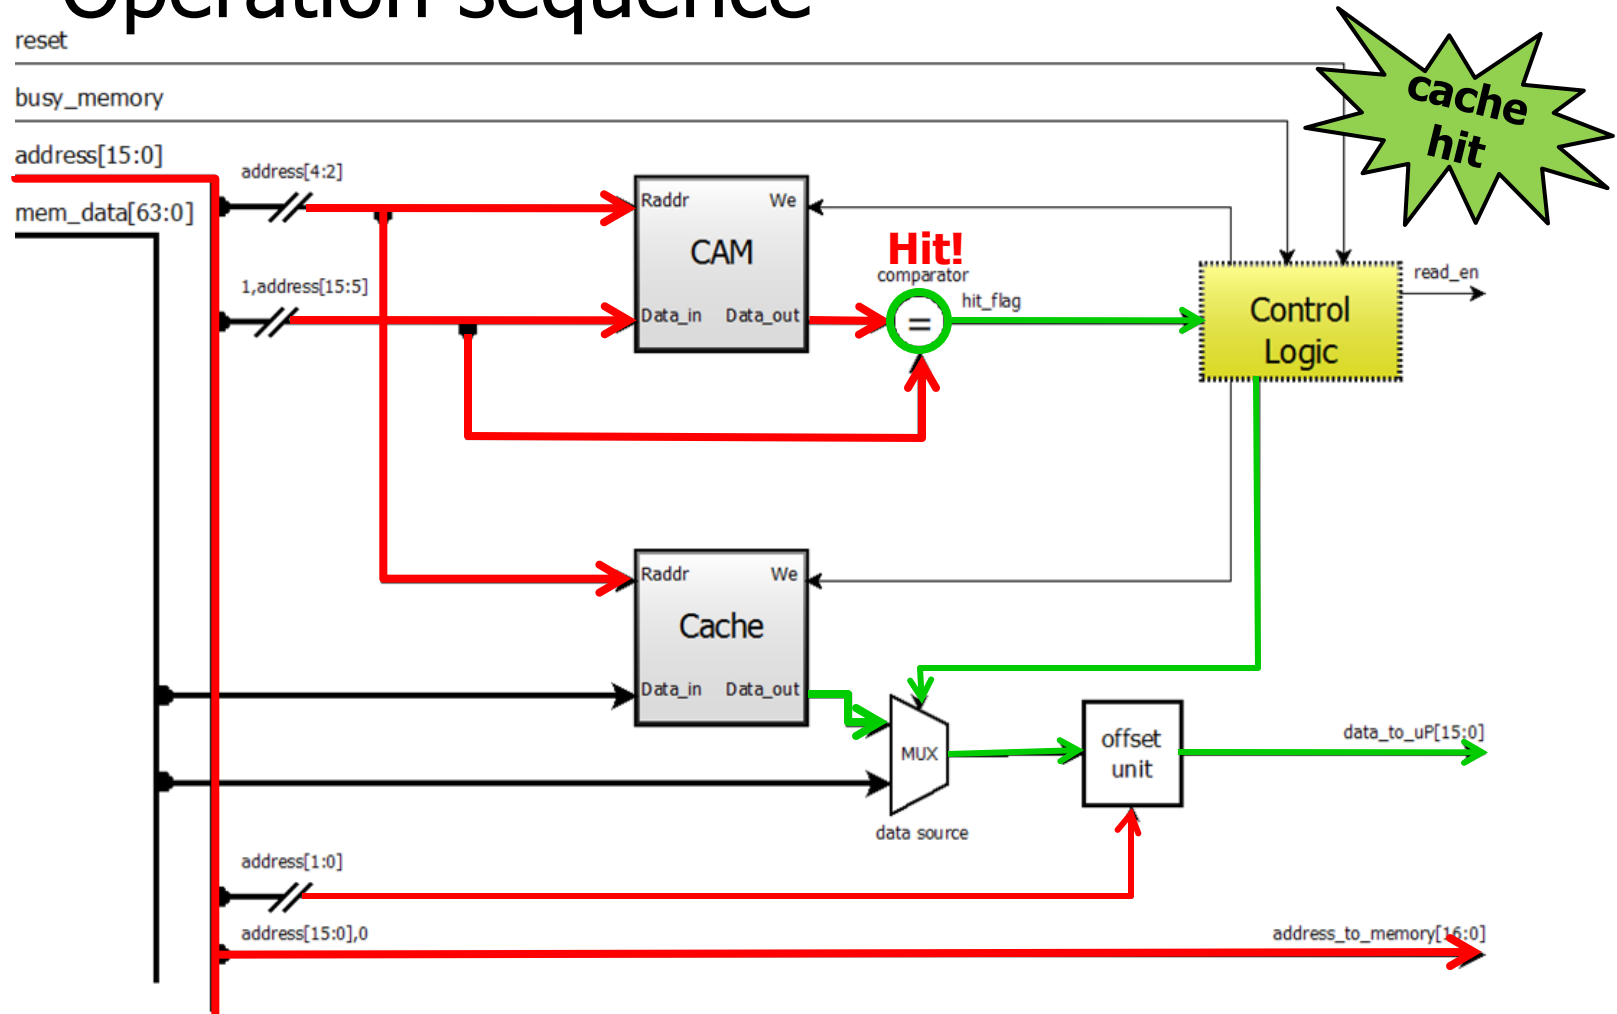
- What about if Cachegen doesn't fit your requirements? (performance, words per line, latency, etc)
- **Lets create a cache from scratch!**
 - **Simplest design: Direct mapped Instruction Cache**
 - What's the main difference between **Instruction** and **Data** caches?
 - **How does it works??**





Working principle basis

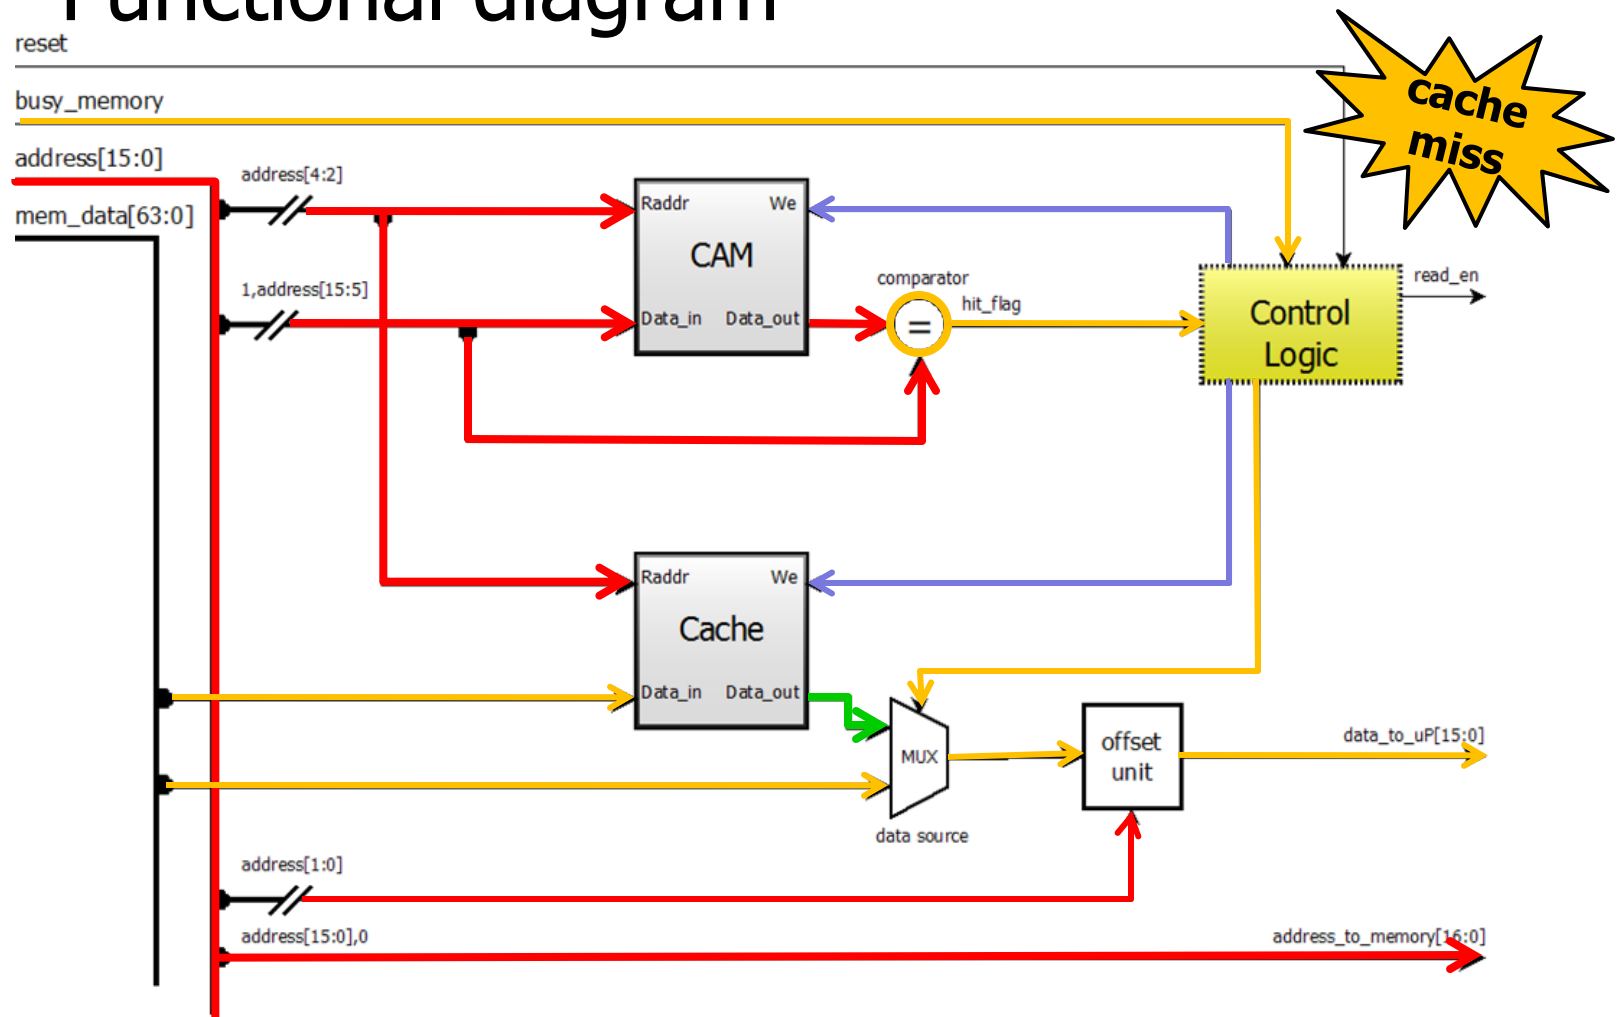
- Operation sequence





Working principle basis

- Functional diagram





Working principle basis

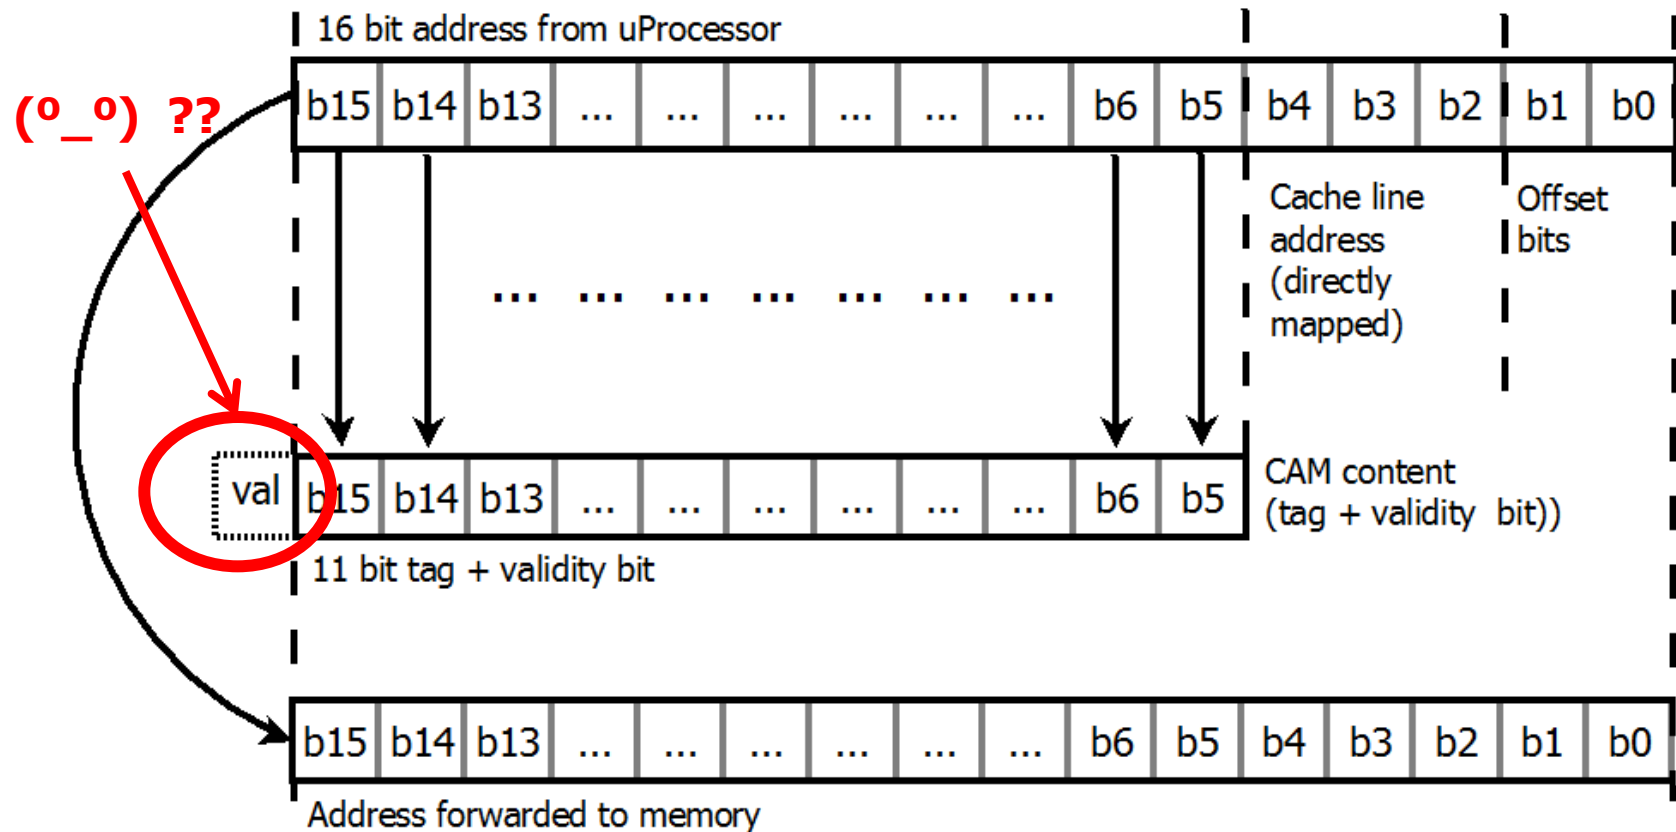
- And what about the **address**? What goes where?
- In multiple words per cache line the **address** must be **split** in
 - Offset , or Word address
 - Line address
 - Tag





Working principle basis

- Address interpretation





Working principle basis

- Address interpretation
 - **Validity bit?**
 - Indicates if a cache line is valid or not.
 - How to know if the cache has stored the address 0x0000 or is just empty? -> **VAL bit**
 - **How many words per line?**
 - **How many cache lines?**
 - **Addressable space?**





Working principle basis

- Now on Vivado;)
 - Create a project
 - Include the "ins_cache.v"
 - Metrics:
 - **How to start...**





Metrics

- SDRAM retrieved big pieces of data
 - Usually memory is organized in long lines 256 bit
 - Retrieved 64 or 72 bits at a time
 - Cache lines should take advantage of that
- Lets assume a **256 byte** memory, with **64 bit lines**
- 32 lines
- Requires **5 bit** to address the lines





Metrics

- Cache set metrics
 - 64 bit memory interface
 - 16 bit architecture
 - 64 Bytes of storage
- 64 bit line
- 64 bit line/16 bit words = 4 Words per line
- $64\text{B}/(4\text{word} \times 2\text{byte}) = 8 \text{ lines}$





- Address Division
 - Address length (PC) 8 bit
 - 4 words per line
 - 8 lines
- 2 bit - word address
- 3 bit - line address
- 8 bit address - (2bit+3bit) = 4 tag bits + val!





Implementation

- Dive into Vivado
 - Analyze and explain
 - Create a Verilog test bench to the project
 - Copy code from “cache_tb”
 - **Run TB!**





Implementation

- How to adapt for **2-way** set associative
 - N-way set associative is basically a direct mapped cache with replicated storage resources
 - N Cache sets
 - N Tag memories
 - N Cache hit comparators
 - ...





Implementation

- How to adapt for **2-way** set associative
 - But where to store the data coming from memory?
 - What's the criteria to select the set?
 - The one which is empty :)
 - And if both addressed lines in all the sets have content?
 - Replacement mechanism must be employed





Implementation

- Replacement Mechanisms
 - Least Frequently Used (LFU)
 - Least Recently Used (LRU)
 - Random
 - Variants of the previous
- **Pseudo LRU**
 - One of the Simplest and Most used
 - Easy to implement





Implementation

- Replacement Algorithm
 - Mechanism used reduce the chances of discard a cache line being used
 - PLRU (Pseudo Least Recently Used) – is a light implementation of the LRU algorithm
 - It has a great efficiency
 - Requires few hardware






- Replacement Algorithm

- **Pseudo LRU**

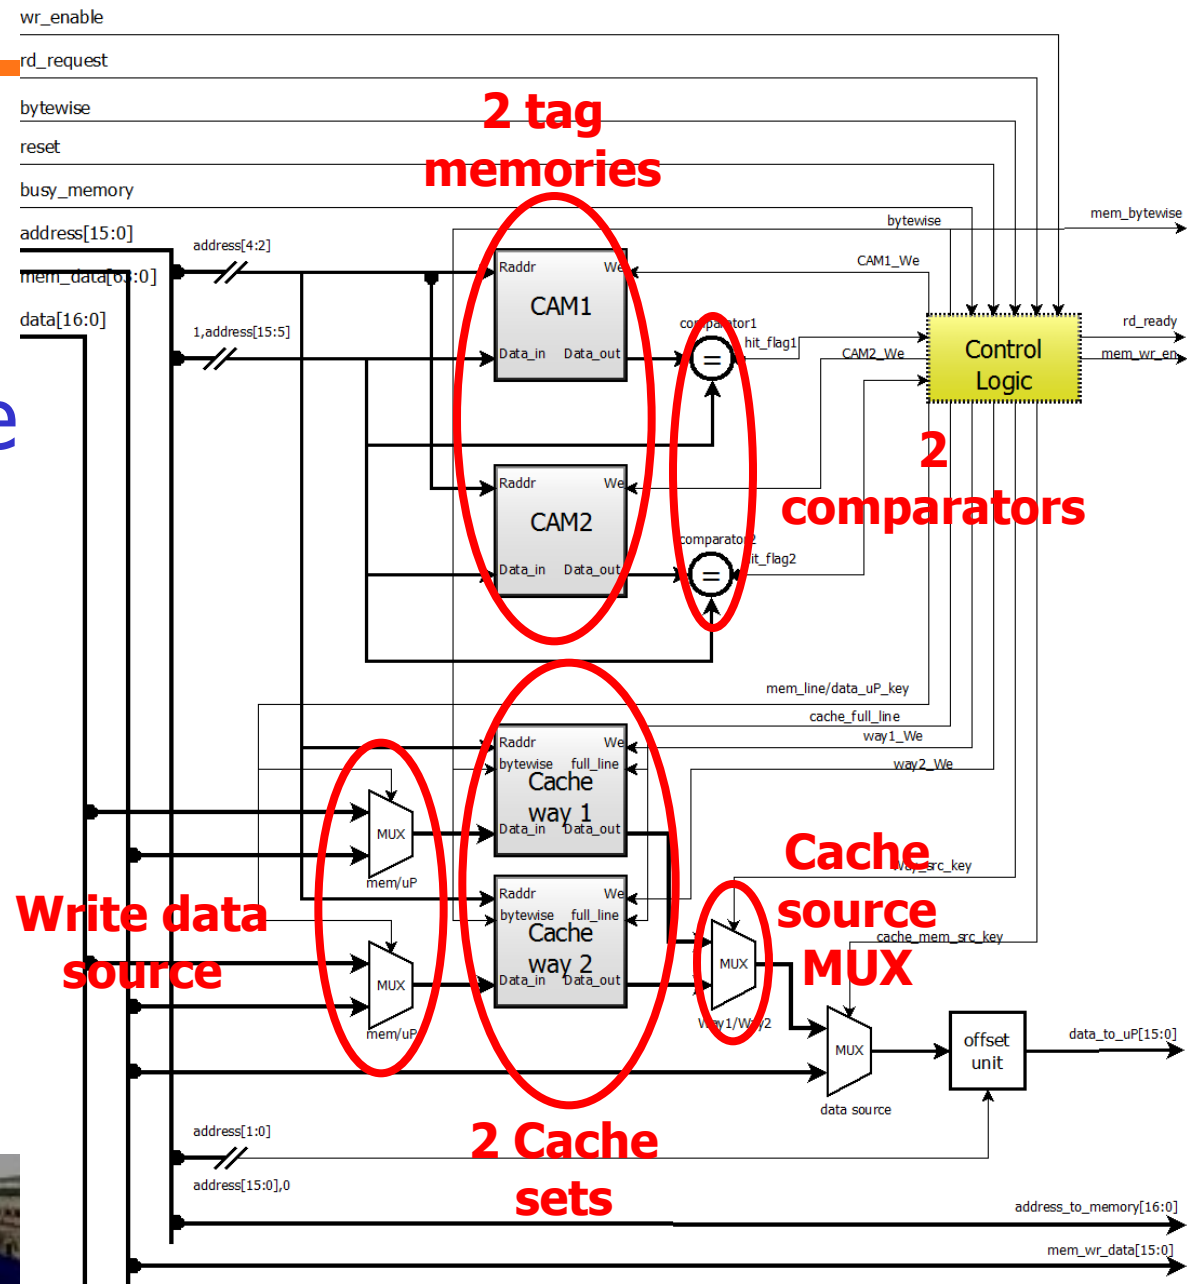
- Uses a binary tree to decide which *way* should be replaced
- Each binary branch (b0, b1, ...) "points the way" to be replaced
 - **0** means "go left"
 - **1** means "go right"
- And "pointer bits" are updated on every access to point to opposite way of the access





Implementation

Two-way
Set assoc.:
Data cache





Implementation

- Try it by yourselves!
 - Create a new Project on Vivado
 - Start from the Direct mapped implementation code
 - Adapt
 - (follow the previous mentioned steps)





Implementation

- Test it
 - Use the provided “mem_hierarchy.v” file, include it in the project
 - It is the Top module that connects the cache and a simulation of an external memory
 - External Memory simulation
 - “Offchip_mem.v” – uses on-chip memory but adds latency cycles
 - Add it to the project





Implementation



Embedded Systems
Research Group

- Thank you for your attention,
- Questions?

